

Experiment 5

Kabir Nawani J043

Sci-kit learn API – Support Vector Classification (SVC)

- **Support vector machines (SVM)** are a set of supervised learning methods used for classification regression and outliers detection.
 - SVMs maximize the margin (Winston terminology: the ‘street’) around the separating hyperplane.
- This becomes a Quadratic programming problem that is easy to solve by standard methods.
- In **Support Vector Classification (SVC)**, the implementation is based on libsvm. The fit time scales at least quadratically with the number of samples and may be impractical beyond tens of thousands of samples.
- For large datasets consider using LinearSVC or SGDClassifier instead, possibly after a Nystroem transformer.

Code:-

```
sklearn.svm.SVC(*, C=1.0, kernel='rbf', degree=3, gamma='scale', coef0=0.0,
shrinking=True, probability=False, tol=0.001, cache_size=200,
class_weight=None, verbose=False, max_iter=-
1, decision_function_shape='ovr', break_ties=False, random_state=None)
```

Parameters:-

- **C(float), default=1.0**

Regularization parameter. The strength of the regularization is inversely proportional to C. Must be strictly positive. The penalty is a squared l2 penalty.

- **kernel{'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'}, default='rbf'**

Specifies the kernel type to be used in the algorithm. It must be one of 'linear', 'poly', 'rbf', 'sigmoid', 'precomputed' or a callable. If none is given, 'rbf' will be used. If a callable is given it is used to pre-compute the kernel matrix from data matrices; that matrix should be an array of shape (n_samples, n_samples).

- **degree(int), default=3**

Degree of the polynomial kernel function ('poly'). Ignored by all other kernels.

- **gamma{'scale', 'auto'} or float, default='scale'**

Kernel coefficient for 'rbf', 'poly' and 'sigmoid'.

1. if gamma='scale' (default) is passed then it uses $1 / (n_features * X.var())$ as value of gamma,
2. if 'auto', uses $1 / n_features$.

- **coef0(float), default=0.0**

Independent term in kernel function. It is only significant in 'poly' and 'sigmoid'.

- **shrinking(bool), default=True**

Whether to use the shrinking heuristic.

- **probability(bool), default=False**

Whether to enable probability estimates. This must be enabled prior to calling fit, will slow down that method as it internally uses 5-fold cross-validation, and predict_proba may be inconsistent with predict.

- **tol(float), default=1e-3**

Tolerance for stopping criterion.

- **cache_size(float), default=200**

Specify the size of the kernel cache (in MB).

- **class_weight(dict) or 'balanced', default=None**

Set the parameter C of class i to $\text{class_weight}[i] * C$ for SVC. If not given, all classes are supposed to have weight one. The “balanced” mode uses the values of y to automatically adjust weights inversely proportional to class frequencies in the input data as $n_samples / (n_classes * \text{np.bincount}(y))$

- **Verbose(bool), default=False**

Enable verbose output. Note that this setting takes advantage of a per-process runtime setting in libsvm that, if enabled, may not work properly in a multithreaded context.

- **max_iter(int), default=-1**

Hard limit on iterations within solver, or -1 for no limit.

- **decision_function_shape{'ovo', 'ovr'}, default='ovr'**

Whether to return a one-vs-rest ('ovr') decision function of shape (n_samples, n_classes) as all other classifiers, or the original one-vs-one ('ovo') decision function of libsvm which has shape (n_samples, n_classes * (n_classes - 1) / 2). However, one-vs-one ('ovo') is always used as multi-class strategy. The parameter is ignored for binary classification.

- **break_ties(bool), default=False**

If true, decision_function_shape='ovr', and number of classes > 2, predict will break ties according to the confidence values of decision_function; otherwise the first class among the tied classes is returned. Please note that breaking ties comes at a relatively high computational cost compared to a simple predict.

- **random_state(int), RandomState instance or None, default=None**

Controls the pseudo random number generation for shuffling the data for probability

estimates. Ignored when probability is False. Pass an int for reproducible output across multiple function calls.

Attribute:-

- ***class_weight_ndarray of shape (n_classes,)***

Multipliers of parameter C for each class. Computed based on the class_weight parameter.

- ***classes_ndarray of shape (n_classes,)***

The classes labels.

- ***coef_ndarray of shape (n_classes * (n_classes - 1) / 2, n_features)***

Weights assigned to the features (coefficients in the primal problem). This is only available in the case of a linear kernel.

coef_ is a readonly property derived from dual_coef_ and support_vectors_.

- ***dual_coef_ndarray of shape (n_classes - 1, n_SV)***

Dual coefficients of the support vector in the decision function multiplied by their targets. For multiclass, coefficient for all 1-vs-1 classifiers. The layout of the coefficients in the multiclass case is somewhat non-trivial.

- ***fit_status_int***

0 if correctly fitted, 1 otherwise (will raise warning)

- ***intercept_ndarray of shape (n_classes * (n_classes - 1) / 2,)***

Constants in decision function.

- ***support_ndarray of shape (n_SV)***

Indices of support vectors.

- ***support_vectors_ndarray of shape (n_SV, n_features)***

Support vectors.

- ***n_support_ndarray of shape (n_classes,), dtype=int32***

Number of support vectors for each class.

- ***probA_ndarray of shape (n_classes * (n_classes - 1) / 2)***

- **probB_***ndarray of shape (n_classes * (n_classes - 1) / 2)*

If probability=True, it corresponds to the parameters learned in Platt scaling to produce probability estimates from decision values. If probability=False, it's an empty array. It uses the logistic function $1/(1 + \exp(\text{decision_value} * \text{probA_} + \text{probB_}))$ where probA_ and probB_ are learned from the dataset.

- **shape_fit_***tuple of int of shape (n_dimensions_of_X,)*

Array dimensions of training vector X